

The background of the slide is a dark gray field filled with a complex network of thin, light gray lines. These lines connect numerous small, dark, spherical nodes. Some nodes are larger and more prominent than others, and the overall pattern suggests a global or interconnected system, possibly representing a neural network or a data network.

Solitaire Using a Deep Q RL Network

Grant Deljevic


The background of the slide is a close-up, slightly blurred image of several playing cards on a green felt surface. Visible cards include the Jack of Spades, Queen of Hearts, and King of Hearts. The text is overlaid on this background.

Solitaire DQN - Introduction


- 3-card Draw Solitaire
- Develop a DQN model to play Solitaire
- Using Reinforcement Learning
- Learn an 'objective' strategy

Solitaire DQN - Deep Q Networks

- DQNs are a class of networks designed for complex state spaces
- Able to make long-term strategic decisions based on card positions and game rules.



Solitaire DQN - Environment: Overview



- Custom built Solitaire environment to simulate the game
- Initially intended to port another environment
- Structure

Solitaire DQN - Environment: Structure

Deck

- Cards
 - Suit (0 – 3)
 - Value (1 – 13)
 - Face up (0 or 1)
- Group Cards
- What can be drawn?

Table

- Cards taken from deck
- 7 Columns
- 4 "top" piles
- Varying cards per column
- Face up or face down

```
def move_possible(self, fromCard, destCard, top=False):
    if not fromCard.face_up:
        return False

    elif isinstance(destCard, int):
        if destCard == -1:
            if fromCard.value == 13:
                return True
            elif fromCard.value == 1 and fromCard.suit == destCard:
                return True

    elif not destCard.face_up or fromCard.value != destCard.value + 1:
        return False

    elif top:
        if fromCard.suit == destCard.suit:
            return True

    elif fromCard.suit % 2 ^ destCard.suit % 2:
        return True

    return False
```

Solitaire
DQN -
Environment:
Possible
Moves

Solitaire DQN - Environment: Functions

Any_legal_moves

- Checks every possible move
- Returns if there's a possible move

Get_possible_actions

- If there are legal moves, creates a list of all possible moves for the network

Check_game_over

- No moves are possible
- Winning or losing conditions are fulfilled

Move_possible

- Checks all game rules for each possible move

Move_card

- Executes the move after checking if it's possible

Set_table

- Sets the initial game state in solitaire format

Environment: Network Interactions

- Get_reward
 - Based on how many cards are face up on the table
 - Large bonus for winning
- Get_state_representation
 - Encode the position of every card
 - Simpler than encoding an action
- Encode_action
 - Encodes possible moves
 - First 2 digits are card's value
 - Next digit is card's suit (red or black)
 - Next digit is the column we're moving a card from
 - Next digit is the column we're moving a card to
 - Next digit is whether or not the destination is the 'top' pile
- Decode_action
 - Decodes action decision
 - Receives values above



Solitaire DQN - Model Architecture

Input of encoded choices & game state

- State size is 26,624

Output of action choice

Training adjusts based on reward

Several dense linear layers

- Hidden layers

ReLU Activation Function

- Allows more complex learning

Final Layer

- Reduces hidden layer to number of possible actions

Solitaire DQN - Policy & Target Network



Policy network acts as the agent and makes decisions, weights are frequently and heavily updated. Learns quickly.



Target network is there to keep a more stable version of weights, and ensure the policy network doesn't collapse.



This allows the policy network to explore aggressively without risk of collapse

Solitaire DQN - Training



Initialize replay memory to store experiences



Training Loop

Reset the environment & receive new initial state for each episode



Process the game state through **get_state_representation** for input into the DQN.



Action Selection (Epsilon-Greedy Strategy):

With probability ϵ , select a random action for exploration.

Otherwise, choose the best action based on the model's prediction (exploitation).

Gradually decrease ϵ over time to shift from exploration to exploitation.



Pass action to environment, receive reward & next state

Solitaire DQN - Training (cont.)



Store the transition (state, action, reward, next state) in replay memory.



Batch Learning

Generate random batch sample

Calculate loss & update the model



Every few steps, update the target network with the weights of the policy network



Log training metrics such as loss and rewards



Assess if the model's performance is converging

Solitaire DQN - Challenges & Solutions



Challenge: Representing the complex state of a Solitaire game

Solution: Develop a function to encode the state into a numerical format understandable by the DQN.



Challenge: Ensure the model chooses valid actions

Solution: Find all possible actions and separately encode them, allowing the network to have the current state while narrowing down the actions it can take to those possible





Challenge: Representing consequences of actions to help the network learn

Solution: Store future states in an action state tree, with probabilities of what revealed cards would be

Note: This is hypothetical

Solitaire DQN - Results and Performance


- Monitoring loss and reward progression.
- Loss was good, reward was not



Solitaire DQN - Conclusions and Learnings



Successfully developed a DQN, though teaching it to play Klondike Solitaire will take more work



Learned about the complexities & pitfalls of applying reinforcement learning to a traditional card game, as well as some of the more intricate parts of DQN

Solitaire DQN - Future Work and Improvements



**Hyperparameter tuning
to optimize performance**



**Finish encoding
of actions and joining it
with the state
representation**



**Creating the action
space tree**



**Optimization – making
functions faster**



Full GPU utilization

Solitaire DQN - Questions?

Solitaire DQN - Q&A

Q1: What is a DQN good for?

- A1: Complex spaces with many options, where a model will need to prioritize long-term strategic thinking

Q2: What was the biggest challenge in training an AI to play Solitaire?

- A2: Creating a quantifiable environment which followed the rules of Solitaire

Q3: How do you measure the success of a trained model?

- A3: Monitor loss, reward, and when that fails, watch it play!

Q4: What is the difference between a policy and target network in DQN?

- A4: The policy network is the more rapidly-updated network which makes action decisions, the target network is there to keep a more stable version of weights.

Q5: What was the ML algorithm used with DQN?

- A5: Reinforcement learning

Solitaire DQN - Citations

1. Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A. A., Veness, J., Bellemare, M. G., & Hassabis, D. (2015). Human-level control through deep reinforcement learning. *Nature*, 518(7540), 529-533.
<https://doi.org/10.1038/nature14236>
2. Sutton, R. S., & Barto, A. G. (2018). Reinforcement Learning: An Introduction (2nd ed.). MIT Press. <http://incompleteideas.net/book/the-book-2nd.html>
3. PyTorch. (n.d.). Documentation. PyTorch.
<https://pytorch.org/docs/stable/index.html>
4. Python Software Foundation. (n.d.). Python 3.9.1 documentation. Python.org. <https://docs.python.org/3/>
5. TensorFlow. (n.d.). TensorBoard: TensorFlow's visualization toolkit.
<https://www.tensorflow.org/tensorboard>
6. Python Software Foundation. (n.d.). Unittest — Unit testing framework. Python 3.9.1 documentation.
<https://docs.python.org/3/library/unittest.html>